

What DeepSeek's breakthrough means for the future of AI (and our portfolio)

Chinese startup DeepSeek shocked the world last week with its new artificial intelligence (AI) model.

It was the first Chinese AI model that matched—and even exceeded—the capabilities of Silicon Valley's finest.

DeepSeek's app shot to #1 on **Apple's (AAPL)** App Store. But that wasn't the reason AI stocks sold off.

The market panicked because DeepSeek supposedly trained its breakthrough model for just \$6 million. That's a fraction of the billions of dollars spent by US companies.

This caused investors to say: "AI just got drastically cheaper. Why do we need all those expensive chips and data centers?"

DeepSeek isn't a threat to the AI boom. It's rocket fuel. This market panic presents a buying opportunity.

Speaking of opportunities, I wanted to quickly remind you of a new project Chris Wood and I have been working on behind the scenes with you in mind.

Our latest venture, *Disruption_X*, was created to pinpoint smaller, "dark horse" disruptors that can 10X or more thanks to the new technologies and industries they participate in.

As part of this project, we studied every US stock in existence over the last 20 years to find every 10-bagger, as well as the common traits that led to them gaining at least 1,000%.

In this month's issue

Special invitation: Join our hunt to find "dark horse" disruptors

DeepSeek's breakthrough

More chips, not fewer

Nothing changes sentiment like price

Who wins? Who loses?

An update on our AI holdings

Key prices

S&P 500	6,037.88
1-Month (Gain) +1.61%	
Nasdaq	21,538.50
1-Month (Gain) +0.10%	
Russell 2000	2,297.42
1-Month (Gain) +0.62%	
Dollar Index	107.95
1-Month (Loss) -0.91%	
US 10-Year Yield	4.511%



A RiskHedge Publication

We wouldn't have been able to pull this project off just a few years ago. It would have required an army of analysts and data scientists trained to look through millions of financial data points.

But thanks to advances in AI, we've programmed a premium large-language model to do the work for us. Then, we analyzed the results ourselves.

What we discovered is a big part of how we'll pick winning stocks in *Disruption_X*.

We put all our findings in our *Ultimate Guide to Modern 10-Baggers* research report, [which you can access here](#). You can think of *Disruption_X* as the natural evolution of *Disruption Investor*, which means you're getting the chance to see it in action first.

We will be launching this new service next week, and as a *Disruption Investor* member, you'll be able to join for a \$500 discount. Be on the lookout for more details very soon.

Now, let's look at the truth behind DeepSeek's breakthrough and what it means for our investments.

DeepSeek's breakthrough

As a result of truly brilliant engineering, DeepSeek's "R1" model is 93% cheaper to run than OpenAI's "o1 Pro," while also matching it in performance benchmarks.

What I love about AI is that you don't have to take anyone's word for it. You can open your laptop and get your hands dirty with it. I tested DeepSeek's model. It's on par with **Alphabet (GOOG)** and OpenAI's best. The big difference is that DeepSeek is free and open source.

But the idea it only cost \$6 million to train is a myth.

In its research paper, DeepSeek clearly states that figure represents only the *final training run* of its R1 model. It excludes all R&D costs... previous experimental runs and testing... data acquisition and processing... and employee salaries and overhead.

Saying DeepSeek built this model for \$6 million is like claiming you can build a house for the cost of hammering in the last nail.

The second myth—that DeepSeek trained its model using cheaper, less-powerful **Nvidia (NVDA)** chips—also crumbles under scrutiny.

My team and I analyzed earlier DeepSeek papers, which referenced having access to at least 10,000 high-end GPUs. Industry experts—including Scale AI's CEO and SemiAnalysis' Dylan Patel—estimate DeepSeek has access to roughly 50,000 premium Nvidia chips.

That means DeepSeek spent upward of \$500 million on its AI infrastructure. Why lie? The US government banned the sale of Nvidia's latest chips to China a few years ago. Admitting to having these cutting-edge chips is admitting to skirting export restrictions.

Our research suggests DeepSeek got its hands on the latest Nvidia GPUs in two ways:

#1: Routing orders through other countries. Nvidia GPU sales in Singapore have surged more than 10X in the past two years. It's highly likely many of these chips ended up in China.


#2: Cloud computing. Sales of physical GPUs to China are banned. But companies across the world can rent GPUs to train their AI systems by simply logging in to US cloud providers like Amazon Web Services. Why try smuggling chips when all you have to do is open your laptop and enter your credit card details?

Let's explore the big investing implications of the DeepSeek fallout.

More chips, not fewer

DeepSeek's R1 model costs 93% less to use than OpenAI's o1. This spooked investors into thinking we'll need fewer chips, data centers, and AI infrastructure going forward.

It's the total opposite. Did you know the cost to use a GPT-4-level model has collapsed 1,000X in the last 18 months...



Source: @thetacosift on X ([click to enlarge](#))

And yet, AI spending has exploded over this period?

Nvidia's AI chip sales surged 600% when AI was getting dramatically cheaper. Why would DeepSeek's efficiency gains be any different?

History shows us each time a technology gets cheaper, it leads to more adoption and spending.

The cost of computing has plunged 99.99999998% since 1990. Did we use fewer chips and hard drives? No, we put computers in everything from toasters to car keys to kids' toys. The same trend played out with cars, electricity, online streaming, and every other innovation I know of. It will also be true for AI.

Cheaper AI means *more* AI everywhere. More applications, more use cases, and more integration into daily life and business. Efficiency gains boost demand for AI infrastructure, not reduce it. The demand for "intelligence" is virtually unlimited.

Expect US AI firms to integrate DeepSeek's efficiency gains into their models soon. Imagine how good AI will be when you take these breakthroughs and scale them up with 10X more compute.

DeepSeek's tech impressed me. But at this early stage when AI is still rapidly improving, efficiency is overrated. The most important thing is creating better AI models. The path toward achieving that remains "more infrastructure."

The big AI spenders won't lift their foot off the gas anytime soon.

DeepSeek showed AI progress is nowhere near hitting a wall, as many feared a few months ago. AI is shifting gears and speeding up. For long-term investors and AI users, that's very exciting.

We've known about DeepSeek since last year when it released its first AI papers. That means the largest AI spenders also knew about it. Two weeks ago, **Microsoft (MSFT)** and **Meta Platforms (META)** both revealed huge AI spending increases.

Would they have made these announcements if they believed DeepSeek's breakthrough would reduce demand? No way.

DeepSeek's breakthrough isn't a setback for AI progress; it's an *acceleration*. More efficient AI will unleash a wave of innovation and adoption. And that's great for AI infrastructure businesses.

Competition is good

The AI race was a high-stake, but friendly, game among US tech kings.

China entering the race changes everything. If America wants to dominate the most important technology of the next 20 years, US companies must go faster.

Top venture capitalist Marc Andreessen called DeepSeek's AI a "Sputnik moment." I agree.

In 1957, the Soviet Union's Sputnik satellite triggered the greatest technological mobilization in US history. Within 12 years, we went from watching communist satellites blink across the sky to placing American boots on the moon.

That same focused determination created Silicon Valley, spawned the internet, and secured decades of American dominance.

DeepSeek is the wake-up call big tech needed in the US. Had this breakthrough come from home soil, the reaction would have been, "Cool, we optimized some algorithms!" But it came from China, which raises the stakes dramatically.

DeepSeek shatters complacency about US AI dominance. It kills arguments for slowing down AI development in the name of safety. It creates the kind of focused competition that's driven America's greatest technological achievements.

AI was already moving at breakneck speed. ChatGPT is a little more than two years old. But it's like a dinosaur compared to the latest AI models. DeepSeek's breakthrough ensures AI progress is about to go hypersonic.

People think DeepSeek means the end of the AI boom. It's the exact opposite. Fierce competition means AI development will accelerate.

Nothing changes sentiment like price

I'm sure you've heard that old Wall Street adage before.

While the long-term case for AI infrastructure is stronger than ever, short-term sentiment has been shaken. Markets were looking for an excuse to sell off, and DeepSeek gave them one.

We expect it will be a volatile few weeks for AI infrastructure companies as markets work through the uncertainty. This volatility is our opportunity to buy great businesses that benefit from faster AI adoption because of reduced costs (more on this below).

But it's important to mentally prepare for more downside in the coming weeks.

One event we're watching that could shift sentiment back is Grok 3's debut.

Elon Musk's xAI is about to release its latest AI model, trained on the largest AI computing cluster ever. "Colossus" contains 100,000 Nvidia H100 GPUs housed in a massive Memphis data center. If Grok 3 is a huge leap forward in capabilities, it's game on. We're also watching the release of OpenAI's new, powerful o3 model closely.

Until then, expect some turbulence. This is a healthy reset for a market that got a little ahead of itself.

Who wins? Who loses?

Before diving into what DeepSeek means for our holdings, let's discuss how it shakes up the AI chess board.

Winners: Big tech

Amazon (AMZN)... Alphabet... Microsoft... and Meta plunged on the DeepSeek news.

Perception: "Big tech spends billions of dollars training AI models while a nimble Chinese startup does it for \$6 million? They're toast!"

Reality: Big tech is the biggest winner from DeepSeek's breakthrough. AI just got a whole lot cheaper. That means big tech can embed it into all their products and services.

Apple was lagging badly in the AI race. Siri is a joke. Apple Intelligence is just ChatGPT wrapped up. DeepSeek changes everything. Future iPhones could run powerful, "free" AI models without draining your battery or sending your data to the cloud.

Amazon didn't create its own leading AI model. Its strategy has been to offer the best tools, regardless of who made them. Amazon can now integrate great, cheap models like DeepSeek into its AWS cloud platform... premium AI capabilities at a fraction of previous costs.

Meta is weaving AI into everything it does, from video recommendations to ad targeting. More efficient AI models allow it to sprinkle intelligence into every interaction, every post, and every ad across its platforms.

Microsoft can continue offering OpenAI's premium models through its cloud platform while also supporting cheaper alternatives. Now developers can access DeepSeek's model for 80% less than OpenAI's offerings through Azure.

Winners: AI users

I'm a heavy AI consumer. Claude is my chatbot of choice. I use it 50+ times a day for everything from editing my writing to teaching my daughter how airplanes fly.

Claude costs \$20/month, which is a steal for how useful it is. DeepSeek means more—and better—AI for free.

DeepSeek was the #1 downloaded app on the App Store last weekend. It caught people's attention, unlike anything I've seen since ChatGPT.

Most folks use free, less powerful versions of AI. DeepSeek was their first taste of using a state-of-the-art model, and I bet they'll want more. AI adoption will accelerate.

DeepSeek is also great for entrepreneurs building AI startups. Until now, you needed deep pockets to access top-tier AI capabilities. DeepSeek changes that equation dramatically.

Entrepreneurs can now leverage world-class AI without needing billions of dollars. A two-person startup can access the same AI horsepower as a tech giant, at a fraction of the cost. Whole new categories of products and services just became possible.

Get ready for an explosion of AI innovation and use cases.

Losers: OpenAI and Anthropic

If OpenAI were a public company, its stock would've likely plunged 30% on the DeepSeek news.

I was about to pay OpenAI \$200/month for its premium “thinking” model, o1 Pro. Then DeepSeek handed you, me, and the rest of the world access to R1 for free.

Unless OpenAI releases a significantly better AI model soon, it could be in big trouble.

Software company ZoomInfo announced that it's replacing OpenAI's models with DeepSeek's R1 for its new AI-powered sales feature. The switch will cut AI costs by two-thirds. Perplexity, the popular AI search engine, is already offering R1 to users.

DeepSeek exposed a big flaw in OpenAI and Anthropic's business models. With clever engineering and model “distillation,” followers can replicate the performance of frontier models at a fraction of the cost.

From Day 1 of the AI boom, we've chosen to invest in the infrastructure layer. The intelligence layer—where OpenAI operates—is ultra-competitive.

Over and over again, OpenAI will debut some awesome new feature that blows the world away. Then, two weeks later, every AI company will have its own version. That's a problem.

Microsoft, which poured billions of dollars into OpenAI, must be asking some hard questions. Why fund massive data centers to train models that competitors can efficiently copy?

We could see Microsoft and OpenAI divorce in 2025. I also expect Anthropic to be acquired, likely by its largest backer, Amazon.

Losers (possible): Memory chip makers

There are two main types of computer chips:

#1: Logic: Nvidia GPUs, which process data.

#2: Memory, which stores data.

Without getting into the weeds, AI needs huge banks of memory. AI servers use 8X more memory than classic computers. And that's doubling with every new chip.

One of DeepSeek's most striking innovations is figuring out how to slash memory requirements by 90% when training its AI model. At first glance, this is bad for memory chip makers like **Micron Technology (MU)** and SK hynix.

But cheaper AI means usage should surge higher. Sure, we'll need fewer memory chips per "query," but there will be so many more queries that memory chip demand could rise.

DeepSeek's efficiency breakthroughs mean AI can now run on your devices instead of through the cloud. Imagine every new smartphone and laptop needing significantly more memory to run AI models "on device." Using less of something in one area often means using much more of it everywhere else.

We're still analyzing the full implications for memory companies and will keep you updated with our latest thoughts.

An update on our AI holdings

Chris here. I want to give you a quick update on our AI stocks considering the turmoil caused by the widespread misinterpretation of DeepSeek's real implications.

Before digging in, let me first say that I'm not recommending selling any of your *Disruption Investor* stocks today.

I also want to say that I recommend against downloading anything from DeepSeek. The terms of service flat out say that your data—and not just data related to the use of DeepSeek—is sent back to China when you use it.

Plus, the software isn't truly open source because you can't verify the training data and full code for security. And as far as I know, there hasn't been any analysis to show one way or the other if running the models "independently" is safe.

But it's your call, of course.

Now, you may recall from last month that I talked about our AI stocks [in the context of the AI value chain](#).

A value chain is simply the set of steps, phases, or stages that add value to a product during its journey of being made.

Understanding value chains helps us identify where different opportunities and risks may lie.

The three broad layers of the AI value chain are:

1. The enabling layer
2. The intelligence layer
3. The application layer

The enabling layer provides the foundation—the bedrock that everything else relies on. In other words, the **“infrastructure”** that makes AI possible. It covers:

- Compute hardware and software like Nvidia’s GPUs (AI chips) and CUDA software platform to fine-tune the chips.
- Memory hardware and software like Micron’s high-bandwidth memory chips and software tools to optimize their performance.
- Networking hardware and software like **Arista Network’s (ANET)** high-performance switches and monitoring software to optimize data flow and prevent congestion and packet loss.
- Power and thermal management solutions like **Vertiv Holdings’ (VRT)** power distribution units and liquid cooling technology.
- Data storage and management tools like **Pure Storage’s (PSTG)** FlashBlade data storage arrays.
- Cloud service platforms like Amazon Web Services and Microsoft Azure.

Without the foundation provided by the enabling layer, AI wouldn’t exist.

The intelligence layer contains the core AI models and algorithms, like the large language models provided by OpenAI (the company behind ChatGPT) and Anthropic (the company behind Claude).

These models provide the “intelligence” businesses can use to build specific applications. Many AI applications use the same underlying models but serve different purposes.

The application layer is where AI interacts with customers in the real world to solve problems and enhance productivity. It includes all the AI-driven products and services that people and businesses use every day.

These solutions leverage AI behind the scenes to generate text and images, automate tasks, and improve customer service. Examples include Anthropic’s Claude, OpenAI’s ChatGPT, Microsoft Copilot, Palantir’s AI software platforms, and **Salesforce’s (CRM)** Einstein AI CRM (customer relationship management) tools.

The AI value chain flows upward through these layers. The enabling layer makes the intelligence layer possible which, in turn, powers the application layer. But the investment dynamics can be quite different at each level:

- The enabling layer often requires massive capital investment but can create strong moats through hardware advantages and scale.

- The intelligence layer typically needs significant research and development investment but can generate valuable intellectual property.
- The application layer typically requires less capital but faces more competition and must focus on user experience and market fit.

Of the 11 AI stocks we currently hold in the *Disruption Investor* portfolio, nine are in the enabling layer, none are in the intelligence layer, and one (Palantir) is in the application layer.

Then there's **Tesla (TSLA)**, which is unique because it's made big investments across all three layers. However, I view Tesla mostly as an application-layer play because its Full Self-Driving-equipped electric vehicles, robotaxis, and Optimus robots bring its AI into the real world to solve problems and enhance productivity for customers.

The DeepSeek story only really affects AI stocks in the enabling and intelligence layers. So, I'll provide a quick update on our nine AI stocks in the enabling layer (because we have none in the intelligence layer).

As you read the updates, please keep in mind what Stephen said earlier about the potential for a few volatile weeks for these companies as markets work through the uncertainty.

Nvidia (NVDA)

As you know, Nvidia dominates the AI chip market. Its GPUs power anywhere from 80%–95% of all AI workloads, according to *Reuters*.

Meanwhile, JP Morgan Asset Management says AI capital spending among the five “hyperscalers”—Alphabet, Amazon, Meta Platforms, Microsoft, and **Oracle (ORCL)**—jumped 56% in 2024 to \$197 billion. And it says this figure will increase to \$234 billion in 2025 and \$249 billion in 2026.

To put that in context, **McDonald's (MCD)** is worth about \$207 billion today... **Shell plc (SHEL)** is worth \$200 billion... and **Disney (DIS)** is worth \$205 billion.

Nvidia GPUs comprise a big chunk of that capital spending in the AI arms race.

Then there's Elon Musk's emerging hyperscaler xAI, which plans to ramp its Colossus AI supercomputer from 100,000 Nvidia GPUs to over 1 million GPUs in the next few years.

In other words, the demand for Nvidia's GPUs remains insatiable. So much so that Elon Musk and Oracle's Larry Ellison were “begging” Nvidia CEO Jensen Huang for more GPUs at a recent dinner they had together.

Meanwhile, as more companies and even sovereign states invest in their own AI clusters, Nvidia's customer base is broadening. This diversification mitigates the risk of over-reliance on a few major customers.

The spending spree on Nvidia's products will slow eventually. But it's not going to happen this year... or next year.

Nvidia's stock peaked at a split-adjusted \$153.13 on January 7, 2025. It dipped as low as \$116.25 in the aftermath of DeepSeek, and it's not trading much above that as I write.

So the stock's off about 24% from its all-time highs. And it looks attractively valued today... especially when you consider the shot of adrenaline the AI market just got from the DeepSeek news. If you don't own a full position, we recommend buying the dip here.

Action to take: Buy Nvidia (NVDA) at current market prices and set a 25% hard stop. (Please note that in addition to the "Free Ride" taken on November 5, 2021, we sold half our remaining shares on February 13, 2024.)

Taiwan Semiconductor (TSM)

Taiwan Semiconductor (TSM) is the "factory" that the world's top tech companies rely on to turn their chip designs into actual physical chips.

It manufactures the vast majority of Nvidia's GPUs and the custom ASICs (AI chips) big tech companies like Alphabet, Microsoft, Amazon, and Meta are developing for AI inference.

On the company's quarterly earnings call held on January 16, 2025, TSM CEO C.C. Wei said:

Even after more than tripling in 2024, we forecast our revenue from AI accelerators [chips] to double in 2025 as a strong surge in AI-related demand continues.

As a key enabler of AI applications, the value of our technology platform is increasing, as customers rely on TSMC to provide the most advanced process and packaging technologies at scale in the most efficient and cost-effective way. To address the structural increase in the long-term market demand profile, TSMC is working closely with our customer to plan our capacity and investing in leading edge specialty and advanced packaging technologies to support their growth.

... Underpinned by our technology leadership and broad customer base, we now forecast the revenue growth from AI accelerators to approach a mid-40% CAGR for the five-year period, starting off the already higher base of 2024. We expect AI accelerators to be the strongest driver of our HPC platform growth and the largest contributor in terms of our overall incremental revenue growth in the next several years.

In other words, demand for TSM's AI chip manufacturing remains insatiable.

TSM's stock hit an all-time high of \$226.40 on January 24, 2025. It fell 17% the following trading day during the DeepSeek panic sell-off but has already regained some of the ground it lost.

Even so, the stock looks like a good buy today. My fair value estimate for TSM is right around \$270 per share, and the stock is trading just under \$210 as I write. So we recommend buying the dip here too if you don't own a full position.

Action to take: Buy Taiwan Semiconductor (TSM) at current market prices and set a 25% hard stop. (Please note: We took a "Free Ride" on October 15, 2024.)

Pure Storage (PSTG)

Pure's flash storage solutions are designed to handle the demanding requirements of AI.

Remember, data fuels AI. It can't continue to move forward without rapidly growing data.

Meanwhile, the hard disk drives (HDDs) the computing industry has relied on to store and retrieve data since the 1960s simply don't cut it anymore. They use too much power and take up too much space.

Pure's flash storage solutions require just one-tenth the amount of power, space, and cooling of HDDs with similar storage capacity. And they produce about one-tenth of the electronic waste. They also require as little as one-fifth the amount of power and just half the space as other flash systems on the market.

Put simply, Pure's better and faster data storage tech allows its customers to save money.

In December 2024, Pure announced that a "top-four hyperscaler"—which means either Amazon, Alphabet, Meta, or Microsoft—chose Pure's DirectFlash technology for its data centers.

Here's some of what CEO Charlie Giancarlo said about the deal on the company's most recent earnings call:

This is the first-ever design win to provide flash for standard hyperscaler storage and it is the vanguard for flash storage, providing all online storage in major hyperscale environments in the future.

... In addition to providing cost-effective data storage, this top four hyperscaler's use of Pure technology is expected to free up significant amounts of power and space in its data centers. It is also expected to significantly reduce the failure rates and maintenance costs associated with legacy disk storage, while doubling the expected lifetime of their storage infrastructure.

... Our design win signals that Pure's DirectFlash technology is now ready to replace hard disks everywhere.

... The design win itself signals that this top four hyperscaler's future data centers are approved to use Pure's technology as its data storage standard.

This is huge news for Pure.

The most important takeaway is that by implementing Pure's tech, big data centers can cut their power consumption by up to 20%. That's game-changing because power is the largest cost for data centers.

Pure's stock hit an all-time high of \$73.67 on January 22, 2025. During the DeepSeek sell-off, it dipped to as low as \$63.44. And while it's already bounced back a bit, we view it as a great buy today.

Please note that our previous guidance was to buy half of your position in Pure Storage up to \$60 per share, and to buy the second half if the stock dips below \$30. Now, our guidance is to buy a full position at current market prices.

So if you already have half a position, now's a good time to pick up the second half. And remember to set a 25% stop based on your cost basis of a full position.

Action to take: Buy Pure Storage (PSTG) at current market prices and set a 25% hard stop.

Synopsys (SNPS)

Synopsys (SNPS) is the leader in electronic design automation (EDA). Chipmakers use EDA software tools to automate semiconductor chip designs. These tools are increasingly critical as chip designs become more complex with the demands of AI.

The company continues to report strong operating results, but the stock has struggled since its highs above \$620 per share in mid-July. That mostly has to do with the fact that Synopsys reduced its backlog and revenue outlook due to US export restrictions on semiconductor technology to China.

China accounted for about 16% of Synopsys' revenue in fiscal 2024. And I don't see the restrictions easing up anytime soon. In fact, we'll probably see more in response to the DeepSeek news.

The good news is that the company's acquisition of **Ansys (ANSS)** is expected to close within the next few months. This will significantly add to Synopsys's competitive advantages in the EDA market and pave the way for faster revenue growth and bigger profits in the future.

See, chip design today consists of 40 or more individual steps performed by a range of software tools. The process includes certain verification steps that all chips must pass before a chipmaker like Taiwan Semiconductor will accept a design for manufacturing.

Ansys provides tools for "power integrity, thermal, and reliability signoff," which are critical verification steps for today's advanced chip designs.

Signoff tools like these are very difficult to replace because they're rooted in years of trust and close collaboration with chip manufacturers.

So, if an EDA company wants to cover the entire chip design flow, it needs Ansys' tech.

Meanwhile, the importance of these signoff technologies is only growing as the market evolves beyond single-chip to 3D (stacked) multi-chip designs. That's because many more physical effects must be simulated in these more complex designs. And that's what Ansys' tools do.

In other words, Synopsys will be even further ahead of its top competitor—and former portfolio holding—**Cadence Design Systems (CDNS)** when this transaction closes in the first half of the year.

So we still see Synopsys as a buy today.

Action to take: Buy Synopsys (SNPS) at current market prices and set a 25% hard stop.

Vertiv Holdings (VRT)

Vertiv provides the critical power management and cooling systems that keep data centers running.

AI computing generates enormous amounts of heat and requires massive amounts of power—much more power than traditional computing. Without Vertiv's solutions, the powerful AI chips would overheat and fail.

Vertiv's stock took a huge hit with the DeepSeek news.

It reached an all-time high of \$155.84 on January 24, then fell more than 35% to a low of \$99.11 the next trading day.

No surprise why. If the DeepSeek news did mean the end of big data center growth, Vertiv's business would be decimated.

But as you know, that's not what it means. So Vertiv's business will be just fine. In fact, the investment thesis for Vertiv is as strong as ever. And I expect the stock to hit a new all-time high within a year.

Action to take: Buy Vertiv Holdings (VRT) at current market prices and set a 25% hard stop. (Please note: We took a "Free Ride" on November 7, 2024.)

Fabrinet (FN)

Fabrinet (FN) is the only company in the world that manufactures the Nvidia-designed high-speed networking cables and optical transceivers that data centers need for AI.

Optical transceivers basically translate one language into another and back again. They convert the digital data computers deal with into light signals—so the data can travel at high speed through the fiber optic cables—and back again.

This relationship with Nvidia has been lucrative for Fabrinet... and the close tie between the two is a big reason why Fabrinet's stock fell 22% on the DeepSeek news.

But there's an emerging trend of unbundling Nvidia's GPU platforms, which makes Fabrinet's continued strong growth less certain.

See, rather than buying fully integrated GPU systems that include additional components, such as Fabrinet-manufactured optical transceivers and interconnects, some big tech companies like Amazon are opting to purchase only the GPUs so they can develop their own customized systems to cut costs.

This emerging trend could significantly impact Fabrinet's business. So we're moving Fabrinet to a "Hold" for now as we work to figure out just how impactful it could be.

Action to take: Move Fabrinet (FN) to a "Hold."

Broadcom (AVGO)

Big tech is now spending billions of dollars developing their own AI chips (custom ASICs) so they're not beholden to Nvidia for AI inference. And they're teaming up with **Broadcom (AVGO)** to do it.

For example, Broadcom has codesigned every generation of Alphabet's custom AI chips called TPUs (Tensor Processing Units).

Broadcom fell more than 17% on the DeepSeek news. It's clawed back some since and, as I write, it's only trading down about 11% from where it was right before the panic sell-off.

We highlighted Broadcom as a top buy less than two months ago in mid-December. It's now about 3% cheaper than it was then. And it looks like a very attractive "Buy" today.

Our original guidance was to wait for a big pullback to pick up the second half of your position. But considering the positive developments that have occurred since we originally recommended the stock in October, we now recommend buying a full position at current market prices.

Action to take: Buy Broadcom (AVGO) at current market prices and set a 25% hard stop.

Monolithic Power Systems (MPWR)

Monolithic Power Systems' (MPWR) chips control and manage the flow of electricity in all sorts of electronics, including Nvidia GPUs.

It's an unsung hero of the AI revolution.

Back in late October, Monolithic sold off after a downgrade from Rosenblatt Securities. It then sold off again in early November when investment firm Edgewater Research released a report—which seemed to be based on a rumor—claiming the company's allocation to Nvidia's Blackwell line of GPUs was "at risk."

We said at the time that no one in our network could confirm the rumor.

To this day, it remains unsubstantiated. There's been no evidence or official statements to support Edgewater's claims.

Monolithic's stock had started to gain some traction in January before the DeepSeek news. On January 23, it closed at \$698.79, which was up 19% from its close just 10 days prior. Then it lost over 11% during the panic sell-off.

We still like Monolithic, but we want a chance to digest the company's fourth-quarter and full-year results that it's releasing after market close on February 6 (today) before we think about moving it back to a "Buy" in the portfolio. So we're keeping it as a "Hold" for now.

Action to take: Continue to hold Monolithic Power Systems (MPWR).

Arista Networks (ANET)

Arista makes high-speed networking hardware and software infrastructure that enables AI by quickly and efficiently moving massive amounts of data among servers, storage devices, and other networked equipment in data centers.

Its customers include the biggest names in AI data centers, such as Microsoft, Amazon, Alphabet, and Meta.

Like several of our other AI stocks, Arista hit an all-time high on Friday, January 24, which was the trading day before the DeepSeek panic sell-off. The stock reached \$133.57, then fell to a low of \$97.68 on Monday, January 27, which reflected a decline of 27%.

Also, like several of our other AI stocks, Arista has bounced back a bit. As I write, it's trading 15% below that all-time high, and the company stands to benefit from the accelerating AI arms race.

Meanwhile, after a couple quarters of revenue growth in the 15%–16% range, we saw year-over-year growth bounce back to 20% last quarter. And on the most recent earnings call, CEO Jayshree Ullal doubled her AI revenue guidance for 2025 to \$1.5 billion.

While the stock doesn't look cheap right now, it *does* look like a good long-term buy at current prices. So we're moving Arista back to a "Buy" in the portfolio.

Action to take: Buy Arista Networks (ANET) at current market prices and set a 25% hard stop. (Please note: We took a "Free Ride" on November 7, 2024.)

Several of our holdings will report earnings before the next *Disruption Investor* publishes on March 6. Here's a rundown on the upcoming earnings releases:

- **Rockwell Automation (ROK):** February 10
- **Vertiv Holdings (VRT):** February 12

- **Taiwan Semiconductor (TSM):** February 12
- **Palo Alto Networks (PANW):** February 13
- **Coinbase Global (COIN):** February 13
- **Arista Networks (ANET):** February 18
- **Zscaler (ZS):** February 24
- **Synopsys (SNPS):** February 26
- **Nvidia (NVDA):** February 26
- **Pure Storage (PSTG):** March 3
- **CrowdStrike Holdings (CRWD):** March 6

Please note the dates listed above are subject to change.

For the most up-to-date guidance on our current positions, [click here](#) to view the portfolio online.

Chris Wood Stephen McBride

Chris Wood and Stephen McBride

Keep up with RiskHedge on the go.

Download the App



Scan it with your Phone

[READ IMPORTANT DISCLOSURES HERE.](#)

YOUR USE OF THESE MATERIALS IS SUBJECT TO THE TERMS OF THESE DISCLOSURES.

Copyright © 2025 RiskHedge. All Rights Reserved