

Nvidia's biggest launch yet... and its impact on our AI holdings

ChatGPT turns three this month.

Since its "birth"—and thanks in no small part to it—artificial intelligence (AI) has gone from fringe idea to the defining tech trend of our time. Its launch was AI's "iPhone moment."

ChatGPT became the fastest-growing app in history, reaching 100 million users 4X faster than TikTok and 35X faster than **Uber (UBER)**.

Today, 800 million folks around the world use it every week to do things like:

- Brainstorm ideas for work and play.
- Generate blogs, social media posts, and other online content.
- Conduct market research.
- Write and debug computer code.
- Build custom AI "assistants" to improve workflows.

I (Chris Wood) work with ChatGPT—and several other AIs—pretty much all day, every day. It's an invaluable teammate that helps make me much more productive.

The muscle behind ChatGPT's growth: **Nvidia (NVDA)**—a *Disruption Investor* Hall of Famer.

When we bought NVDA in September 2020, it was a \$320 billion market cap chip company. **Late last week, it became the world's first \$5 trillion company.**

In this month's issue

Happy birthday, ChatGPT!

A review of the AI value chain

Big tech is still spending billions on AI

The ripple effects of Nvidia's "rack scale" revolution

Key prices

S&P 500	6,771.55
1-Month (Gain)	+0.46%
Nasdaq	23,348.64
1-Month (Gain)	+1.77%
Russell 2000	2,427.34
1-Month (Loss)	-2.37%
Dollar Index	100.22
1-Month (Gain)	+2.56%
US 10-Year Yield	4.085%



A RiskHedge Publication

What's crazy is that Nvidia hit this \$5 trillion milestone less than three months after breaching the \$4 trillion mark.

To put that trillion-dollar jump in value into perspective, consider that **Berkshire Hathaway (BRK.B)** is worth about \$1 trillion today. It took super-investor Warren Buffett about 60 years to build Berkshire into a trillion-dollar company. And Nvidia added that much value virtually overnight.

Meanwhile, Nvidia is now worth about \$1 trillion more than **Apple (AAPL)** and \$1.1 trillion more than **Microsoft (MSFT)**, the second- and third-largest companies in the world by market cap, respectively.

We were early to the AI infrastructure buildout, taking profits on NVDA twice as it exploded higher. We're currently sitting on a 602% gain, meaning every dollar we've invested has grown 6X.

We still think Nvidia will perform extremely well in the coming years. In fact, I think it will become a \$10 trillion company before 2030... possibly in 2028. I also think most investors are missing important aspects of Nvidia's story, which significantly bolster the case for our other AI holdings.

I'll tell you all about that in just a bit. For now, suffice it to say we're happy holding what's left of our Nvidia position.

Meanwhile, as we've written about in the past two issues of *Disruption Investor*, we're entering Phase 2 of the AI boom.

This is where the bottleneck in compute is giving way to bigger bottlenecks in power, networking, cooling, and memory—and where compelling opportunities in the “application layer” of the AI value chain are starting to emerge, setting the stage for a whole new set of winners.

As Disruption Investors, we want to go where the money is headed—not where it's already been.

So today, we're giving you a full update on our “enabling layer” AI holdings.

We believe these stocks are still best-positioned to win as they continue to lay the groundwork for a new wave of “AI native” disruptors in the “application layer” coming down the pike.

A review of the AI value chain

In our [January](#) and [February](#) 2025 issues of *Disruption Investor*, we talked about our AI stocks in the context of the AI value chain.

A value chain is simply the set of steps, phases, or stages that add value to a product during its journey of being made.

Understanding value chains helps us identify where different opportunities and risks may lie.

The three broad layers of the AI value chain are:

1. The enabling layer
2. The intelligence layer
3. The application layer

The **enabling layer** provides the foundation—the bedrock that everything else relies on. In other words, the “building blocks” or “**infrastructure**” that makes AI possible. It includes things like compute, memory, storage, networking, power, and cooling. Without the foundation provided by the enabling layer, AI wouldn’t exist.

The **intelligence layer** contains the core AI models and algorithms—like the large language models provided by OpenAI (the company behind ChatGPT) and Anthropic (the company behind Claude)—companies can use to build specific applications. Many AI applications use the same underlying models but serve different purposes.

The intelligence layer also includes AI frameworks like **Alphabet’s (GOOGL)** TensorFlow and **Meta Platforms’ (META)** PyTorch. Developers use these tools to help them build the AI models.

The **application layer** is where AI meets the real world. It includes all the AI-driven products and services that people and businesses use every day to solve problems and enhance productivity.

The AI value chain flows upward through these layers. The enabling layer makes the intelligence layer possible, and both power the application layer.

Most of our AI investments so far have fallen into the enabling layer.

That makes sense because we’re still in the early days of AI’s story. (But stay tuned for the opportunities coming in the application layer, where more “AI native” companies are starting to look attractive.)

Our enabling layer AI holdings include: Nvidia, **Taiwan Semiconductor (TSM)**, **Micron Technology (MU)**, and our Phase 2 AI “ETF” consisting of **Vertiv Holdings (VRT)**, **Arista Networks (ANET)**, and **Broadcom (AVGO)**.

Before we dig into the company-specific updates, I want to provide some color on why our enabling layer stocks will continue to benefit from the AI infrastructure boom... and where the next wave of AI spending is heading.

Big tech is still spending billions on AI

ChatGPT was AI’s first killer app. But it won’t be the last.

Just like the internet gave us **Amazon (AMZN)**, Alphabet, Uber, and many other disruptors, AI will drive the creation of a new wave of businesses that couldn’t exist without this new infrastructure.

Today, most of these companies are small, pre-earnings, and far from household names. These are the sort of stocks we track in our ***Disruption X*** advisory.

In August, for example, we recommended the second AI-native stock that passed our screens, which could become the first \$1 trillion business of its kind. You can learn more [here](#).

Our *Disruption Investor* playbook remains the same: Follow the money. Own the AI stocks getting paid today—while taking some profits and letting them run—as we hunt down the next generation of winners.

Remember, the main reason our AI stocks *are* getting paid today is because of big tech—especially the big four US-based hyperscalers: Meta, Alphabet, Microsoft, and Amazon.

It’s their money—their capital expenditures, or “CapEx”—that’s driving this whole AI infrastructure buildout boom. So it’s their money we want to follow first.

These four companies just reported quarterly results last week, so we have fresh data and management commentary to gauge how things are progressing.

I crunched all the numbers, and what I found bodes extremely well for our AI stocks.

Here’s the story...

You may recall that we were a bit concerned about big tech CapEx earlier in the year.

On the *RiskHedge Reserve* call in mid-April, I said:

I think in the very near term, because of how globally exposed these big tech companies are—and how cyclically exposed they are—a slower-moving economy could (and probably will) result in a very near-term pullback in AI CapEx.

In this current environment of uncertainty, and due to the fact that it takes time for supply chains to adjust, these companies are going to want to defend their margins. And that probably means a deceleration in CapEx...

I still couldn't be more bullish on the investment opportunities for AI in the long run. But the next several months warrant some caution.

Q1 '25 numbers came out a couple weeks after that call, and it looked like our concern was warranted.

CapEx for these four hyperscalers totaled \$76.6 billion in the first quarter. That was up more than 60% on a year-over-year basis, but down about 2% sequentially from Q4 '24—marking the first quarterly CapEx decline in eight quarters.

Management teams mostly cited supply-side issues like lengthy power purchase agreement negotiations, securing grid hookups, and permitting delays as the reasons for the decline. But there was good reason for investors to be concerned about the demand side of things going into Q2.

What Q2 and Q3 told us, however, is that our concern was misplaced.

All four hyperscalers significantly increased their AI investments (both on a year-over-year and sequential basis) in the second and third quarters of the year. And all four increased their CapEx guidance in Q2 '25 and again in Q3 '25.

Their combined CapEx in Q2 was \$95.0 billion, up 24% from Q1. And their combined fiscal-year 2025 CapEx guidance increased 15% from \$306 billion to \$352 billion.

Their Q3 combined CapEx (just reported last week) jumped to a whopping \$112.5 billion. That's an annual run rate of \$450 billion.

And based on what the CFOs of these companies said on their earnings calls last week, we know their combined CapEx will exceed that \$450 billion over the next year.

Here's some specific CapEx commentary from those calls:

- Alphabet CFO Anat Ashkenazi: "We now expect CapEx to be in the range of \$91 billion to \$93 billion in 2025, up from our previous estimate of \$85 billion... Looking out to 2026, we expect a significant increase in CapEx."
- Meta Platforms CFO Susan Li: "Our current expectation is that CapEx dollar growth will be notably larger in 2026 than 2025."

- Amazon CFO Brian Olsavsky: “Looking ahead, we expect our full-year cash CapEx to be approximately \$125 billion in 2025, and we expect that amount will increase in 2026.”
- Microsoft CFO Amy Hood: “Next, capital expenditures. With accelerating demand and a growing RPO balance [backlog], we’re increasing our spend on GPUs and CPUs. Therefore, total spend will increase sequentially, and we now expect the FY ‘26 growth rate to be higher than FY ‘25.”

Digging into that Microsoft comment: The company’s fiscal 2025 CapEx was \$64.55 billion, and that was up 45.1% from fiscal 2024. That means the company expects FY ‘26 CapEx to be at least \$94 billion.

Meanwhile, we’ve recently seen a flurry of similarly aggressive AI infrastructure investment outside of these four hyperscalers.

For example, OpenAI, emerging hyperscaler **Oracle Corp. (ORCL)**, and Softbank announced plans to spend an additional \$300 billion building five new US-based AI data centers as part of their Stargate Project. That’s in addition to their prior \$100 billion commitment.

And don’t forget about Elon Musk. His xAI—an emerging hyperscaler like Oracle—is hard at work building Colossus 2, which will initially incorporate over 500,000 GPUs, and plans to scale to millions.

Long story short: This surge in AI infrastructure investment is a strong tailwind for AI stocks. And there’s no need to worry about big tech CapEx for at least the next year (though we’ll continue to monitor it).

The reason this investment will continue unabated for the foreseeable future: These companies simply can’t afford to fall behind in the AI race.

When you believe—as they do—that the ultimate prize is a “digital-god-scale” platform worth trillions of dollars over decades, you’re highly motivated to keep investing as much and as fast as possible.

As you’ll see in the next section, falling behind is exactly what would happen if they weren’t spending as aggressively as possible on Nvidia’s new AI systems.

The ripple effects of Nvidia's "rack scale" revolution

With hyperscaler spending ramping up, our enabling layer AI stocks are entering a critical moment. Let me walk you through what's happening across the portfolio—and how we're managing each position today.

Nvidia (NVDA)

I'm going to stop calling Nvidia a "chip company" because it's evolved into so much more.

Today, Nvidia is an "AI systems company."

Its premiere AI system is the GB300 NVL72. This is a game-changing product. I think it—along with the GB200 NVL72, which launched a few months earlier—is arguably the most significant product launch in Nvidia's history.

This product is also something most investors appear to be overlooking because it marks the shift from 8-GPU "server scale" systems to 72-GPU "rack scale" systems, which has far-reaching implications across the AI infrastructure ecosystem.

A server is basically a single metal computer box—like a beefed-up PC with no display screen or keyboard—packed with components like GPUs, CPUs, and memory chips. A rack is a six- to seven-foot-tall metal cabinet that holds multiple servers stacked vertically like pizza boxes.

"Server scale" refers to building systems by focusing on what's inside one server box. For example, Nvidia's older setups—like the DGX H100 AI computer—houses eight GPUs linked together to work as one large GPU. If you want more power, you connect multiple servers together with standard external networking.

"Rack scale" means building the entire rack (the whole tall shelf) as one giant, unified AI system. Nvidia's new GB300 NVL72 packs 72 Blackwell Ultra B300 GPUs (plus 36 Grace CPUs, high-bandwidth memory, interconnects, etc.) into a single rack, linked with lightning-fast networking to act as one giant GPU.

The key difference is that server scale is modular and slower (relatively speaking), while rack scale is integrated and faster.

This shift to rack-scale systems is game-changing for Nvidia because it's evolving from selling what are more like AI "kits" to turnkey integrated AI machines. This lets Nvidia:

- Control more of the design (like networking, cooling, and software).
- Charge premium prices (a single GB200 NVL72 rack sells for about \$3 million and the newer GB300 system sells for about \$3.5 million to \$4.5 million depending on the configuration).
- Push the boundaries on performance when it comes to training and running massive AI models with 4X faster training, 30X faster inference, and 50X higher AI factory output per megawatt.

For big tech customers like Microsoft and Alphabet—as well as “neocloud” providers like **CoreWeave (CRWV)** and AI “startups” like OpenAI building data centers—it's a big win for speed and efficiency.

It's also a much simpler setup because you don't have to fiddle with tons of cables between servers. And it lets you operate a smaller footprint or get more out of your current footprint because one rack can do what used to take many racks.

Put simply, you haven't seen anything yet when it comes to Nvidia's hardware capabilities.

Nvidia is shipping GB200 NVL72 systems at scale and just started shipping GB300 NVL72 systems. Customers with the cash and need for industrial-scale AI are gobbling them up as fast as they can... because if they don't, they'll get left in the dust.

In the company's fiscal Q2 '26 (April 28, 2025, through July 27, 2025) “data center” revenue—which closely approximates AI revenue—grew 56% year over year to \$41.1 billion.

That's an annual run rate of \$164.4 billion. But as big as that number is, it's just the beginning.

Nvidia CEO Jensen Huang recently said he's chasing a \$3 trillion to \$4 trillion AI infrastructure opportunity over the next five years... and that he has “visibility” into \$500 billion in data center revenue for the six quarters through the end of calendar year 2026.

Huang wasn't saying Nvidia would capture that entire \$500 billion chunk over the next six quarters, but my projections aren't far off that mark.

I like to lean to the conservative side when making financial forecasts to bake in some cushion. Even so, my model has Nvidia generating about \$370 billion in data center revenue over the next six quarters.

And for the three calendar years of 2026 through 2028, I have Nvidia producing data center revenue of about \$270 billion, \$400 billion, and \$600 billion, respectively.

Nvidia currently trades at a price-to-sales multiple of about 30. That sounds extremely rich, but it's not considering how fast the company is growing.

If I take a more conservative estimate of 17 (to reflect the fact that growth will slow over time) and multiply that by my \$600 billion data center revenue forecast for calendar year 2028, we get \$10.2 trillion.

In other words, I think Nvidia could be a \$10 trillion market cap company by the end of 2028.

That number is so big, it's hard to wrap your head around. But I think it's completely reasonable.

The stock won't go up in a straight line from here, and we could certainly see a correction of around 40% (or more) within the next few years. But that will only be a buying opportunity if things stay on track from an operational and demand perspective.

Action to take: Buy Nvidia (NVDA) at current market prices and set a 60% hard stop. (Note: We took a "Free Ride" on November 5, 2021. We then sold half our remaining shares on February 13, 2024.)

As I alluded to earlier, Nvidia's shift to a rack-scale AI systems provider significantly bolsters the bull case for our other enabling layer AI holdings.

That's because when you change the "unit" of computing from a server to an entire rack, everything around that rack—the networking, power, cooling, memory, and other supporting components—must scale up and adapt.

You can think about it like upgrading from a 200-horsepower engine you might find in a Toyota Camry to a 1,500-horsepower Bugatti engine. You can't just drop the new engine into your Camry because it would tear apart the chassis and drivetrain, melt the brakes, shred the tires, and cause complete structural and mechanical system failure almost instantly.

An upgrade like this requires a complete ecosystem overhaul. Shifting from server-scale to rack-scale systems requires a similar overhaul.

In other words, Nvidia's move to rack-scale AI systems is *not* a solo act. It's a spark that ignites a virtuous cycle for its "pit crew" of supporters and suppliers.

Now, I'll walk you through how each of our enabling layer AI holdings stands to benefit from it.

Arista Networks (ANET)

Arista provides the super-fast “data highways” for AI setups.

It makes networking switches and software that connect racks together in massive AI clusters. If each GB200 or GB300 NVL72 rack is like a city of 72 GPUs, it is Arista’s highways that let these cities communicate with each other at incredible speeds.

Let me explain.

Training and running advanced AI models requires thousands of GPUs working together across many racks. These GPUs must constantly talk to each other—sharing data, results, and updates with minimal delay. Plus, some of these GPUs need to use outputs from others as their inputs.

So if the network is slow or congested, those expensive GPUs sit idle waiting for data, which drags down performance. In other words, even the most powerful GPUs are held back if they’re connected to a lousy network. This has made networking as important as the chips themselves.

Nvidia’s rack-scale AI systems underscore this point.

Inside an NVL72 rack, Nvidia uses its proprietary NVLink technology to connect the 72 GPUs with extremely high bandwidth—creating an internal GPU network far faster than traditional connections. So fast, it’s like one giant GPU.

But between racks—like when connecting multiple NVL72 racks together to scale beyond 72 GPUs or linking to storage systems—data must still travel over external network switches. So these external links need to be just as fast and scalable to avoid becoming the bottleneck.

Arista specializes in exactly that kind of high-speed, low-latency networking tech that large AI deployments require.

Meanwhile, when you treat racks as the unit of compute, you need to move much more data between racks as fast and reliably as possible. That drives demand for faster, smarter switches and software to coordinate traffic across an AI cluster. And that’s where Arista excels.

Most new rack-scale AI systems Nvidia ships will likely be accompanied by orders for Arista’s superfast switches to connect that rack to the world. So Arista becomes the key supplier of the “data highways” Nvidia’s multimillion-dollar NVL72 systems need to reach their full potential.

I must note that Nvidia is becoming more of a competitor to Arista these days with its new Spectrum-X Ethernet networking stack. But demand for this technology will far outstrip supply in the coming years, so there’s plenty of room for multiple players to thrive.

What's more, due to Arista's open standards, broader compatibility, proven deployments, and deeply entrenched relationships with hyperscalers like Microsoft and Meta, it's clear that the shift to rack-scale AI systems represents an enormous opportunity for Arista.

Action to take: Buy one-third of a full position in Arista Networks (ANET) at current market prices and set a 40% hard stop.

Note: Arista is part of our Phase 2 AI "ETF," along with Broadcom and Vertiv Holdings. That's why we only recommend buying one-third of a position. We're treating all three companies together as one full position.

Broadcom (AVGO)

Broadcom provides the "traffic control center" networking chips that direct data flow in AI setups.

These chips basically work as the "brains" inside the switches provided by Arista (and others like Cisco and Juniper) to connect racks together in AI clusters.

Rack-scale AI means more switches running at higher speeds with advanced features to handle AI workloads. Broadcom is the dominant chip supplier in this space, so whenever companies like Arista sell more switches, Broadcom benefits through increased networking chip sales.

In fact, Broadcom holds about 90% market share in Ethernet switch chips. So the company benefits essentially every time a GB200 or GB300 NVL72 rack gets connected in the construction of an AI cluster.

If you scale a cluster like this to 100,000 GPUs, you'd need a few thousand (probably 3,000 to 5,000) Broadcom switch chips. (It's tough to say exactly how many because the scale-out network needs between racks isn't disclosed in public documents.)

Meanwhile, Broadcom is making strategic moves to tailor its products to emerging networking challenges so it can remain on top in the age of rack-scale AI.

Just last month (October 2025), for example, Broadcom launched "Thor Ultra"—a cutting-edge switch chip built to handle the massive cluster sizes AI demands. According to the company, the goal of the new chip is to entrench its "control of network communications inside data centers designed for AI."

Broadcom also benefits from rack-scale AI by designing custom AI accelerator chips for big tech customers like Alphabet and Meta.

The gist of the story here is that big tech companies don't want to be so reliant on Nvidia and its GPUs forever—at least on the inference side of AI. Plus, you can gain efficiencies in speed and cost over general-purpose GPUs when you build custom AI accelerators for specific sorts of workloads.

So big tech is now spending billions developing their own AI chips. And they're teaming up with Broadcom to do it.

This leads to huge benefits like:

- **Big contracts with recurring revenue:** When Broadcom designs a chip just for Alphabet, Meta, or some other big tech company, it locks in years of orders. And just a single deal can be worth billions.
- **Higher revenue per chip:** Custom chips are complex and unique, so Broadcom can charge premium prices—way more than off-the-shelf chips.
- **Full-package sales:** In rack-scale AI, everything (chips, networking, and cooling) is designed together. So Broadcom can incorporate its other products—like switch chips, network interface cards, PCI Express links, and optical transceivers—into these designs, allowing it to generate much more revenue on a per-customer and per-deal basis.
- **Stickiness:** When a customer builds their whole AI system around Broadcom's tech, switching to a competitor is extremely expensive and slow. So Broadcom becomes entrenched in the future of these big tech companies.

Long story short, while Broadcom is not as visibly associated with AI as Nvidia, it's one of the big “quiet” winners of the AI boom. Nvidia's rack-scale systems supercharge the need for robust networking, and Broadcom's tech is at the heart of many networking solutions.

Action to take: Buy one-third of a full position in Broadcom (AVGO) at current market prices and set a 40% hard stop.

Note: Broadcom is part of our Phase 2 AI “ETF,” along with Arista Networks and Vertiv Holdings. We're treating all three holdings as one full position.

Vertiv Holdings (VRT)

Vertiv provides the power management and cooling tech for AI setups.

Vertiv's power management tech ensures electricity is available where and when it's needed in a data center. It keeps servers, cooling systems, storage systems, and other critical equipment up and running.

And the company's cooling tech keeps all that critical equipment from overheating.

In other words, Vertiv's products ensure racks of servers get all the electricity they need and that the tremendous heat they produce is safely removed.

Rack-scale AI significantly ups the ante for Vertiv's tech.

See, back in 2020, average server rack power densities in data centers were typically in the range of 5 to 9 kilowatts (kW) per rack. Racks were loosely filled and ran warm but not dangerously hot. So big room air conditioners blowing cold air through the aisles could handle data center cooling needs.

But AI changed all that.

A single Nvidia DGX server with 8 GPUs (previous generation "server-scale") could draw around 5 kW, which amped up rack power densities to the 30-40 kW range.

Nvidia's new rack-scale NVL72 systems take power requirements to another level.

According to **Hewlett Packard Enterprise (HPE)**, which builds racks based on Nvidia's designs, each GB200 NVL72 rack can consume about 132 kW of electricity under full load. GB300 racks suck even more power.

That means a single rack now draws more power than 100 average American homes. And it's a *dense* concentration of power. So each rack generates an enormous amount of heat. Way too much to be cooled by traditional air-based systems.

That's where Vertiv's liquid-cooling tech comes in.

Liquid cooling leverages the superior thermal transfer properties of water and other fluids to cool high-density servers up to 3,000X more effectively than air.

It involves either circulating a coolant like water through cold plates in direct contact with electronic components like GPUs or submerging electronic components in a non-conductive liquid coolant.

Vertiv has a full portfolio of liquid-cooling solutions to meet the thermal management needs of rack-scale AI today and tomorrow.

Vertiv's technology is so important for rack-scale AI that Nvidia chose to partner with the company when designing the new GB200 and GB300 NVL72 systems. Specifically, Vertiv co-developed reference designs for direct-to-chip liquid cooling and high-voltage DC power management.

While that doesn't mean that Vertiv's hardware is necessarily embedded in every NVL72 rack that sells, it *does* mean to me that the company's products will be the de facto default solutions here.

We're seeing this in real-world deployments. CoreWeave's first cloud-based GB300 NVL72 system uses Vertiv's CoolChip in-rack coolant distribution units and PowerDirect DC shelves to handle about 140 kW per rack.

Put simply, Vertiv's tech is integral to rack-scale AI and only becomes more important as power and cooling needs continue to grow.

As AI clusters scale, Vertiv's infrastructure doesn't just sustain them—it allows Nvidia's chips, Broadcom's chips, and Arista's switches to operate at peak performance, turning raw compute into trillion-parameter breakthroughs.

Action to take: Buy one-third of a full position in Vertiv Holdings (VRT) at current market prices and set a 40% hard stop.

Note: Vertiv is part of our Phase 2 AI "ETF," along with Arista Networks and Broadcom. We're treating all three holdings as one full position.

Micron Technology (MU)

Micron makes the high-bandwidth memory (HBM) that feeds GPUs the massive amounts of data they need to access in real-time to handle AI workloads.

Imagine trying to do research but only being able to see a few words at a time before needing to flip pages, or pull another book off the shelf to see just a few more words. Your research would take forever.

Now imagine being able to see (and make sense of) thousands of pages from countless sources all at once. You can work infinitely faster because all the information you need is right in front of you. That's like the difference between traditional memory technology and HBM for AI.

Micron is one of three companies in the world (in addition to SK Hynix and Samsung) capable of providing cutting-edge HBM at scale. And it's a huge beneficiary of Nvidia's computing revolution because the more GPUs and CPUs in a rack—a defining feature of rack-scale AI—the more memory the rack needs.

According to Micron CEO Sanjay Mehrotra, "Micron's HBM3E 8-high is designed into Nvidia's GB200 system, and our HBM3E 12-high is designed into the GB300."

That means Micron's HBM is validated and used in some GB200 and GB300 NVL72 systems. Nvidia sources HBM from all three vendors—about 60% from SK Hynix, 25% from Micron, and 15% from Samsung today—to ensure supply chain diversity and meet massive demand.

So roughly a quarter of the NVL72 racks Nvidia sells today ship with Micron's HBM.

I must note that both SK Hynix and Micron have already effectively sold out their 2026 HBM supply. In other words, all the HBM the two companies can produce next year is already spoken for at agreed prices.

Meanwhile, Micron is the sole supplier of what's called LPDDR5X memory for the 36 Grace CPUs in NVL72 racks. So every rack that ships *does* contain Micron's tech. This CPU memory is separate from GPU memory but is still vitally important for managing overall system operations and performance.

What all this means is that Micron has essentially hitched a ride on Nvidia's success. When Nvidia sells a rack-scale AI system, Micron is selling memory inside that system.

Action to take: Buy Micron Technology (MU) at current market prices and set a 60% hard stop.

Taiwan Semiconductor (TSM)

Taiwan Semiconductor—or TSMC, for short—builds the “brains” (GPUs, CPUs, and custom ASICs) for AI setups.

But that's not all...

A rack-scale AI system like the GB200 or GB300 NVL72 is an orchestra of advanced chips: dozens of GPUs, dozens of CPUs, stacks of memory, and networking chips all work in concert in each rack.

TSMC is the primary manufacturer of almost all those advanced chips.

So every time Nvidia sells one of these racks, it's stuffed with high-end silicon made by TSMC.

TSMC is also the sole provider of advanced packaging technology needed for HBM. As AI scales and needs to incorporate more and more memory, TSMC's packaging expertise becomes more and more valuable.

The company also benefits from the premium pricing and complexity of all these products. So as Nvidia shifts to ever-more advanced chips and integrated systems, TSMC earns more from both wafer fabrication and packaging per unit sold.

It's a great relationship: Nvidia gets to push the envelope with record-breaking AI systems, and TSMC reaps the rewards by being the exclusive manufacturer enabling that push.

Action to take: Buy Taiwan Semiconductor (TSM) at current market prices and set a 40% hard stop. (Note: We took a “Free Ride” on October 15, 2024.)

Two of our holdings will report earnings before the next *Disruption Investor* publishes on December 4:

- **Palo Alto Networks (PANW):** November 19
- **Nvidia (NVDA):** November 19

Please note that the dates listed above are subject to change.

And for the most up-to-date guidance on our current positions, [click here](#) to view the portfolio online.

Chris Wood Stephen McBride

Chris Wood and Stephen McBride

[READ IMPORTANT DISCLOSURES HERE.](#)

YOUR USE OF THESE MATERIALS IS SUBJECT TO THE TERMS OF THESE DISCLOSURES.

Copyright © 2025 RiskHedge. All Rights Reserved