

The hidden hero “feeding the AI beast”

“How do you use AI?”

That’s my new go-to question when meeting contacts.

Admittedly, it’s self-serving. I get to pick the brains of some of the smartest investors and thinkers in the world about how they’re using the most powerful technology of our lifetimes.

Artificial intelligence (AI) chatbots have transformed my work life. I use Claude, Grok, and ChatGPT at least 50 times daily as my research assistant, writing partner, and instant answer machine. It’s like having a brilliant intern who never sleeps, never complains, and is always ready to help.

I’ve even taught my kids how to use AI. Now, when they ask me something I don’t know, they immediately say, “Ask the lady on your phone”—their nickname for ChatGPT’s voice mode. That simple phrase reveals so much about how this technology is becoming woven into our daily lives.

I’ve always liked technology you can get your hands on and understand right away.

What separates AI from other megatrends is its immediate accessibility. You don’t need special equipment or technical expertise to test a world-class AI—just open your laptop, click a few buttons, and maybe pay \$20/month.

One major reason we founded *Disruption_X* this year is because we can see what AI is rapidly enabling. This is a tool of extreme leverage that lets small teams punch far above their weight class.

In this month’s issue

How does AI improve?

The next stage of the AI boom

Will there be one AI to rule them all?


Feeding the AIs

The hidden hero powering the AI revolution

INOD’s path to profits

America’s fastest-growing stocks

The data backs this up. Look at this chart from Stripe showing how fast AI startups are growing:



Source: Stripe

Never before in history have we seen startups scale as fast as today's AI companies. Not during the internet boom, not during the mobile boom, not ever. Let that sink in.

I laugh when people call AI a fad. Those who think this technology won't reshape our world simply haven't been paying attention. Or they tried the original ChatGPT two years ago, found it mediocre, and haven't bothered with the quantum leaps made since.

Two years ago, I barely used AI. It was fine for creating silly pictures or asking ChatGPT to write a limerick about Rhodesian Ridgebacks, but it wasn't truly useful for serious work.


Now? The landscape has completely transformed. Almost monthly, a cutting-edge AI model debuts that pushes the frontier forward. There are almost too many breakthrough models to keep track of.

Take xAI's Grok 3, which launched just weeks ago. I've used it daily since its release, and it's mind-blowing. You can feed it 100,000 words of messy notes and ask it to analyze specific patterns or insights, and it delivers value instantly—a task that would have taken a human analyst days or weeks.

This rapid progress isn't slowing down. Look at this graph showing how AI models have been steadily conquering benchmarks, with OpenAI's o3 and Grok 3 leading the pack:

Frontier Model Intelligence Over Time by AI Lab

Artificial Analysis Intelligence Index includes MMLU Pro, GPQA Diamond, Humanity's Last Exam, LiveCodeBench, SciCode, MATH-500, AIME 2024
Intelligence Index estimated via interpolation for certain models



Source: Artificial Analysis

What's driving this improvement? That's where this month's investment story begins.

How does AI improve?

That's the \$300 billion question, which is roughly how much **Amazon (AMZN)**, **Microsoft (MSFT)**, **Alphabet (GOOG)**, **Meta Platforms (META)**, and **Oracle (ORCL)** will spend building AI data centers this year alone.

It's not magic. It's math—and it follows rules called **scaling laws**.

The two raw ingredients you need to make AI better are: data (books, essays, and Wikipedia pages) and compute (**Nvidia [NVDA]** GPUs).

Scaling laws are brutally simple: Crank up the data and computing power by 10X, and the AI's quality roughly doubles.

This is why Elon Musk's xAI built a monster cluster of **100,000 GPUs**, triple the size of the previous record. More GPUs processing more data creates better AI.


The new Grok 3 was trained with **3e26 FLOPs** of compute. It's equivalent to running your smartphone continuously for **634,000 years**. That astronomical computing power is how Grok 3 smashed benchmarks in record time.

At the same time, the amount of data used to train top AI models has increased by a mind-numbing **1,018,766,656%** since 2010.

It's hard to wrap your head around "1 billion percent." It's like turning \$1 into \$10 million. That's the kind of exponential growth AI has achieved in just over a decade.

In 1950, Claude Shannon built one of the earliest AI examples, a robotic mouse named Theseus that "remembered" its path through a maze. Theseus learned from exactly 40 data points (the walls it bumped into).

Today's leading AI models train on **trillions** of data points, as this chart shows:



Source: Our World in Data

Today's AI models are essentially pattern recognition systems on steroids. They identify and learn from text, images, and other data. The more quality data they access, the more they grasp the nuances and complexities of our world.

But if your data is trash, your AI is trash, no matter how sophisticated your model.

GPT-2 (which came out in 2019) trained on roughly 3 billion words. The entire English Wikipedia, for comparison, contains about 4.6 billion words.

Just four years later, GPT-4 trained on approximately **9.75 trillion words**—over 2,000 times the entire English Wikipedia.

This isn't some weird academic "who has the bigger AI" contest. Data is the difference between AI being a cool toy and a transformative tool.

The original ChatGPT (3.5) scored a mediocre 30% on GPQA (a challenging AI benchmark). The data-enhanced GPT-4.5 hit 70%.

More data = better AI.

The next stage of the AI boom


I know many *Disruption_X* members joined us from our big brother, *Disruption Investor*. And one of our biggest wins there was being early on Nvidia.

We rode Nvidia's GPU empire from the foothills to the summit. It was a beautiful journey.

But here's a crucial insight about technology waves: **The most obvious plays get crowded fast.** The infrastructure boom has been the headline story so far. Everyone from Wall Street analysts to your **Uber (UBER)** driver knows about it.

While spending on this physical AI buildout will continue climbing as the tech giants race to build their Digital Gods, what matters to us as investors is the "rate of change."

In plain English: The percent growth is slowing down, even as the absolute dollar amounts hit record highs.



Source: @modestproposal1 on X


Does this mean AI's boom is fizzling? No. It's the exact opposite.

The hardware race was just the opening act. Now, the spotlight shifts toward **AI services**—the companies that feed, deploy, and harness these hungry new machines.

I know we're still early because **less than 5% of people worldwide** use AI tools like ChatGPT today. Compare that to smartphones, where over 90% of Americans now carry one. But it took 17 years to get there.

AI is on the same adoption curve. Five years from now, I bet half the planet will use AI daily. That's a massive runway for growth, which drives revenue to innovation.

This chart is more evidence we're early:



This isn't a chart of OpenAI or Google's AI sales. It's consulting giant **Accenture's (ACN) AI-related revenues**, which hit a staggering **\$3 billion** in 2024 and continue to climb.

How is a consulting firm making billions from AI when most companies are completely lost?

A parade of Fortune 500 firms are desperately trying to figure out how to use this technology before their competitors do. But most haven't even completed step one: getting their data organized, cleaned, and ready to feed into AI systems.

That's what they're paying Accenture billions to help them figure out.

Accenture doesn't even do the heavy lifting itself. It assess the situation, makes recommendations, and then points these companies toward specialized data experts who can actually clean up the mess.

This is how raw and early the market is. It's like companies in 1998 hiring consultants to help them figure out how to set up an email server.

Will there be one AI to rule them all?

For the past few years, the AI story has been dominated by the race to build colossal, all-purpose models. Digital equivalents of Swiss Army knives that can handle everything from writing poetry to coding software.

That's because the bigger the model, the better its answers.

Financial news outlet *Bloomberg* created "BloombergGPT" to gain an edge in analysis and forecasting. It fed the model every scrap of data *Bloomberg* had gathered over 40+ years.

GPT-4, which wasn't specifically trained for finance, smoked it in tasks like figuring out the sentiment of earnings calls. It didn't matter how much high-quality, "proprietary" data *Bloomberg* had because GPT-4 was 100X bigger.

It was the same story when Google's general-purpose AI model went head to head with a model specifically trained for healthcare. Google's general model won handily.

Lesson: If you want the best AI, just make it massive.

But there's a major shift underway, one that savvy investors must understand. The future isn't one AI to rule them all. It's a flourishing ecosystem of specialized, on-premises models optimized for specific jobs.

Three powerful forces are driving this shift:

1. **Keep your crown jewels safe.** If you're running a bank, hospital, or defense contractor, you simply cannot send your most sensitive data to GPT-4 or Claude. Full stop. Even for companies without regulatory constraints, the risks of exposing proprietary data are too great.
Would **Coca-Cola Co. (KO)** send its secret formula to an external AI? Not a chance. These companies will bring the AI in-house where their data never leaves the security perimeter.
2. **Call an expert.** China's AI model DeepSeek handed us a gift when it pioneered the "Mixture of Experts" (MoE) approach. Instead of one massive model where every parameter works on every problem, MoE divides the model into specialized "experts." When you ask a question, the system routes it only to the relevant experts. For example, the math expert for calculations or the reasoning expert for logic problems. Since only a fraction of the model activates for any given task, you get faster responses, lower costs, and better performance.
Every AI company is now adopting MoE.

3. **A unique edge.** The brutal truth about AI models is that you can train them on as many GPUs as you want, but without unique data, they're commodities racing to the bottom. Open-source models are catching up to proprietary ones at breakneck speed.

The only sustainable advantage is **unique data no one else can access**. This is why I'm bullish on:

- xAI's Grok, fueled by X's social network;
- Meta's models, trained on Facebook, Instagram, and WhatsApp;
- Google's Suite, which can leverage the vast treasure troves of YouTube and Gmail.

This is also a huge advantage for specialized, on-premises models built on proprietary company data.

The next wave isn't building one AI to rule them all. It's building thousands of specialized AIs, each ruling its own domain.

Companies building specific AIs on their unique data assets will create billions of dollars in value. These models are being crafted right now, largely out of the public eye.

Feeding the AIs

After investing in dozens of megatrends over my career, I've learned one crucial lesson: **You have to be nimble**. After identifying a trend early and riding the initial wave, it's easy to get caught in the hype and overstay your welcome.

We were early to the AI data center buildout because few spotted its importance. Now, that story is on CNBC daily. When Wall Street catches up to your thesis, it's time to look for what's next.

The best opportunities usually come from small, nimble sectors where the crowd isn't looking yet. Everyone's obsessed with flashy AI models and the chips that power them.

But they're missing a critical question: **What goes into these models before they're even built?**

Data is the lifeblood of AI. No data, no AI. It's that simple.

Feeding data into AI models isn't a simple copy-paste job. According to S&P Global, a staggering **60% of an AI developer's time** isn't spent building cool chatbots or features. It's spent on the grueling work of gathering, preparing, and cleaning data.

I'm sure you've heard the analogy "data is the new oil." It's imperfect, as data can be reused infinitely and grows rather than depletes. But in one crucial respect, the comparison is spot-on: **Both must be refined before they're valuable.**

When crude oil comes out of the ground, it's a toxic, unusable sludge. It must go through complex refinement processes before becoming the gas that powers your car.

Data is the same. The raw information companies collect is messy, inconsistent, and often incomprehensible to AI systems.

A global retailer like **Target (TGT)** has reams of customer data: purchase histories, browsing patterns, demographic information, in-store traffic... all scattered across hundreds of systems in dozens of formats.

Refining it into machine-readable data that can be fed into AIs is the real moneymaker.

Companies that can clean and structure this data won't just participate in the AI revolution, they'll lead it. And this trend of refining data becomes even more important as we move into the world of unstructured data.

Text is structured and it's mostly been tapped, at least what's available on the open web. But video and audio remain a huge pool of untapped potential. Think of all the YouTube videos, movies, documentaries, and songs ever recorded. This is the motherlode that remains largely untouched.

The catch is that this media is a blob until it's been cleaned, labeled, and made machine-readable. That's the bottleneck... and our big opportunity.

While everyone else fights over scraps in crowded spaces, the companies that can effectively feed the AI beast with quality data will quietly build the foundation of trillion-dollar empires.

I'll hand you off to Chris to tell you about our newest recommendation doing just that.

The hidden hero powering the AI revolution

"Garbage in, garbage out"...

I'm sure you're familiar with the concept.

If you eat bad food, your body's not going to work like it's supposed to.

If you cook with bad ingredients, your food's not going to taste good.

If you base your business decisions on bad information, your products or services are going to fail... and so on.

"GIGO," as it's often called in the tech world, is also a core principle of computer science. It applies in the same way here. The quality of what you put into a system determines the quality of what you get out of it.

GIGO is particularly important when it comes to AI.

Let me explain...

Traditional software systems follow explicit, human-written rules. If the programmer writes buggy code (inputs bad data), the software won't work as intended. But the result (output) will be poor in a predictable, one-to-one way.

AI systems, on the other hand, learn patterns from their input data and then apply those patterns over and over again. This creates an amplification effect where small input data issues can lead to big, broad problems.

Think about a traditional database that contains one incorrect record—one piece of bad input data. Only queries about that specific record will return the wrong information. But if an AI model trains on that same incorrect record, it might incorporate the error into its understanding of an entire concept, affecting countless future outputs.

What's more, in traditional software, we can inspect the code to see exactly how inputs become outputs. The process is transparent and follows human-designed logic. AI systems like neural networks, on the other hand, develop their own internal representations even their creators can't see. This "black box" nature means:

- It's difficult to detect when garbage inputs cause garbage outputs.
- It's tough to trace how specific inputs influence specific outputs.
- It's challenging to fix problematic patterns without disrupting beneficial ones.

There are other issues like feedback loops and emergent behaviors that can compound AI's GIGO problems, but you get the idea.

GIGO highlights the fact that despite their impressive capabilities, AI systems remain fundamentally dependent on the quality of the input data.

Data quality is crucial in artificial intelligence because it directly impacts AI models' accuracy, performance, and reliability. High-quality data is the cornerstone of any successful AI.

Think of it like the high-octane fuel that powers a supercar. If you start with lower-grade fuel, even the best car with the best engine won't perform to its potential. And if you keep feeding it that lower-grade fuel, you'll eventually break it.

When you feed AI systems high-quality data—that's accurate, consistent, complete, timely, representative of the real world, and relevant to the problem(s) at hand—they:

- Learn accurate patterns rather than misleading correlations;
- Make reliable predictions that work well even in new situations;
- Avoid perpetuating harmful biases;

- Require less time and computational resources to reach optimal performance;
- Foster trust and confidence among users.

Put simply, high-quality data isn't just "nice to have." It's the very foundation on which trustworthy AI is built. And as AI's influence spreads, the importance of high-quality data only grows.

The problem is that virtually all organizations—even the biggest tech companies with all the money in the world—lack the expertise needed to properly prepare data for AI consumption.

Let me walk you through an example of what it takes to generate high-quality data for AI to illustrate this point...

In this example, you're a healthcare company developing an AI to help identify early signs of lung cancer in chest X-rays.

Step 1: Data Collection & Curation

First, you need to collect chest X-rays spanning years from thousands of patients who were diagnosed with lung cancer, and thousands who weren't. The goal is to gather a diverse set of images that represent different patient demographics, imaging conditions, and stages of lung health.

Then, you organize the images systematically. This involves categorizing the X-rays based on factors like age, gender, and medical history, and ensuring that each image is accompanied by relevant clinical information like patient outcomes. This organization makes it possible to manage and utilize the data effectively.

The expertise needed for this Data Collection & Curation includes:

- Radiologists to identify relevant cases;
- Medical researchers to design a representative dataset;
- Legal/privacy experts to ensure patient data is handled properly;
- Data engineers to extract data from various hospital systems, clinics, and imaging databases;
- IT experts to design storage and access systems.

Step 2: Data Cleaning & Preprocessing

Raw data can be messy and is never ready for immediate use. Cleaning it up ensures the AI can "read" it clearly.

So your team reviews the X-rays and removes or fixes any that are blurry or have technical artifacts. Clear, high-quality images are vital for accurate analysis.

Then, you standardize the images, so they have a consistent format, size, brightness, and contrast. These preparations help the AI process the images effectively.

The expertise needed for this Data Cleaning & Preprocessing includes:

- Data scientists who know how to normalize imaging data;
- Medical professionals who can identify technical artifacts in X-rays.

Step 3: Annotation & Labeling

The AI needs to know exactly what lung cancer looks like on an X-ray. This is where you “teach” it what to look for. It’s like showing a child pictures of dogs and cats and saying, “This is a dog, and this is a cat.”

Experienced radiologists examine each X-ray and mark areas of concern that look like they could be lung cancer (called nodules). They label these regions with tags like “suspicious nodule” or “normal tissue.”

The expertise needed for this Annotation & Labeling includes:

- Software developers to create annotation tools;
- Senior radiologists to create annotation guidelines;
- Multiple radiologists to perform annotations (with protocols for disagreements);
- Project managers to track annotation progress and quality.

Step 4: Data Validation & Quality Assurance

This is like double-checking your work to make sure everything is correct and consistent.

It’s important because even the best doctors and data scientists make mistakes, and you need to ensure the data is accurate, so the AI doesn’t learn from errors.

This step might involve having a second team of radiologists review a random sampling of the annotations to confirm findings and implementing protocols to regularly check the quality of both the images and their annotations to ensure the data remains reliable over time.

The expertise needed for this Data Validation & Quality Assurance includes:

- Quality assurance specialists with both medical and data science knowledge;
- Statisticians to design sampling protocols;
- Senior medical experts to serve as final arbiters.

Step 5: Data Enrichment & Augmentation

This step is about enhancing your data set to improve the AI's learning experience.

You may enrich the data by adding information to each X-ray, such as patient symptoms, smoking history, or genetic factors to provide more context for the AI system.

To increase the diversity of your dataset without collecting new X-rays, you may apply techniques like rotating or flipping existing images. This helps the AI become more robust in recognizing lung cancer signs under various conditions.

It's like practicing a skill in different ways—reading a book, listening to an audiobook, and watching a movie about it. Each version helps better your understanding.

The expertise needed for this Data Enrichment & Augmentation includes:

- AI researchers who understand data augmentation techniques;
- Medical experts who can verify augmented data remains realistic.

Step 6: Documentation & Governance

This is like keeping a detailed diary of everything you've done. It helps others understand your process and ensures you're not breaking any laws or ethical guidelines.

For each X-ray, your records must show where it came from, who annotated it, how it was cleaned and processed, and any other modifications made.

You also need to establish policies to manage access to the data to ensure compliance with medical regulations, such as patient privacy laws.

The expertise needed for this Documentation & Governance includes:

- Data documentation specialists;
- Data governance experts;
- Compliance officers.

The Goldengate advantage

As you can see, each of the steps outlined above requires significant expertise from an array of experts across different fields. There's no way your healthcare company could have all this expertise.

And as you can imagine, each of these steps also requires significant time. In fact, according to AI research firm Cognilytica (now part of Project Management Institute), developing high-quality data takes up 80% of the time for most AI projects.

Meanwhile, even though the example I mentioned above was for a healthcare company trying to improve lung cancer detection—the expertise issue I’ve been talking about applies to every other company, too.

The expertise issue becomes even more substantial when you’re trying to build bigger, broader AI systems like the multimodal—meaning capabilities across text, images, audio, and video—Large Language Models (LLMs) that power OpenAI’s ChatGPT, Anthropic’s Claude, xAI’s Grok, etc.

As I mentioned earlier, even trillion-dollar big tech companies lack the expertise necessary to generate the high-quality data needed for AI...

That’s why they’re partnering with a small \$1.3 billion market cap company called **Innodata (INOD)** for their AI data needs.

Innodata was founded 35 years ago to “engineer the highest-quality data so organizations across broad industry segments could make smarter decisions.”

Today, that means creating the “teaching materials” for AI.

The company has over 6,600 employees and associates across 31 countries supporting more than 40 languages to help create the high-quality data (teaching materials) AI learns from.

Many of these folks hold advanced degrees in fields like law, business, finance, medicine, and other hard sciences, social sciences, and technology.

Innodata doesn’t just create the teaching materials for AI manually. The company uses its proprietary AI platform called Goldengate—which it built and refined over the past decade—to make the process much more efficient and scalable.

Goldengate is an AI system that helps build AI systems. It takes care of the initial processing of raw data. Human experts review and correct the AI’s work. These corrections teach the AI how to improve.

As the AI improves, it can handle more complex tasks with less human intervention. This frees up human experts to focus on more sophisticated aspects of data preparation.

It’s a virtuous cycle.

What’s impressive is Goldengate can ingest and make sense of raw “unstructured data” like videos and audio recordings. This was impossible for machines to do just a few short years ago.

I spent a decade trying to find an investable company with technology that can make sense of unstructured data because it accounts for 80% to 90% of the world’s data, in terms of both volume and value.

Tech giants building advanced AIs and smaller companies across all industries just trying to improve their businesses need to be able to leverage this kind of data. But it was an untappable resource until recently.

Goldengate isn't the only system that can make sense of unstructured data. But its ability to ingest this data—and help turn it into the structured teaching materials AI needs to learn from—makes it special.

Also, Goldengate is “low-code,” meaning you don't need a team of data scientists to use it. Subject matter experts without programming skills can work with the platform to create high-quality training data for AI.

To get a better sense of how things work, imagine Innodata is preparing financial documents to train an AI to analyze loan agreements.

Goldengate would first scan thousands of loan documents and automatically identify and annotate key information like interest rates, terms, collateral requirements, etc. Financial experts would then review these automated annotations, correcting any errors and adding nuanced interpretations that the AI missed.

Goldengate would record these corrections and use them to improve its understanding. For instance, if experts consistently correct how it interprets certain financial clauses, it would learn these patterns. Over time, it would get better at recognizing complex financial concepts automatically, requiring less human intervention.

What makes Innodata's approach effective is that it doesn't try to fully automate the process. The company recognizes that creating the high-quality data its customers need requires both:

1. AI systems that can process massive amounts of information very quickly;
2. Human experts with specialized knowledge who understand nuance and context.

By combining these strengths, Innodata can create better teaching materials than either humans or AI could produce alone. And it can create the enormous volumes of high-quality data big tech companies need to build advanced AI systems.

We've seen this proven by the fact that eight of the world's largest tech companies (including five of the “Magnificent 7”) now rely on Innodata for their AI data needs.

While Innodata can't name its specific big tech customers publicly, a few I think I've been able to figure out are **Apple (AAPL)**, Microsoft, Alphabet, Amazon, Nvidia, and OpenAI.

To give you a small sense of some of the problems Innodata is solving for its big tech customers, let's look at a real-world case study:

A “leading technology company” needed a specialized dataset to train its AI (an LLM chatbot) to perform complex “multi-action chaining” tasks.

This would allow the AI to tackle much more complex problems and do things like deep research and scheduled tasks. It works by teaching the AI to break up big jobs into smaller tasks, to execute each action in sequence, and to use the output from one action as the input for the next.

That may sound simple to us, but it's quite a challenge to teach an LLM—which relies solely on patterns learned during initial training—new tricks.

Innodata's team of experts cleverly engineered prompts that coaxed the AI into revealing its internal thought process while responding to users.

Using the insights from this exercise, Innodata's team of skilled authors crafted a dataset that consisted of thousands of simulated chatbot conversations. These conversations were able to teach the AI to perform complex, multi-action tasks and significantly improve its practical applications in the process.

This sort of “supervised fine-tuning” of AI models—which Innodata excels at—is something all the LLM developers need. It's how you transform these language models into systems that can follow specific instructions and perform useful tasks on their own.

Innodata's revenue growth & big opportunity

In the fourth quarter of 2024, Innodata's revenue grew 127% year over year to \$59.2 million. And revenue for the full year jumped 96% from the previous year to \$170.5 million.

And I see Innodata's solutions becoming even more important as the AI market develops. This bodes extremely well for future growth, which is what matters most.

As the AI market develops, three important trends we're seeing evolve are:

- AI models becoming more complex;
- Inference becoming a bigger focus;
- Synthetic data becoming more important.

As AIs become more complex “reasoning” models, the need for Innodata's services grows because the need for high-quality data grows.

Advanced AI models can learn more nuanced patterns and relationships and solve more complex multi-action chain problems, but only if they learn from the highest-quality data.

You can think of it like how an intelligent college student can learn much more advanced concepts than a kid in grade school, but he or she needs to be provided with appropriately sophisticated educational materials.

What's more, as AI systems become more powerful, ensuring they behave safely and align with human values becomes increasingly critical. This requires specialized data for alignment, safety testing, and evaluation—all areas where Innodata excels.

Meanwhile, as these complex AIs are increasingly adapted for specialized domains like healthcare, finance, and legal applications, they need more domain-specific data and expertise. This increases the value of Innodata's subject matter experts across various fields.

Now, on to the “inference becoming a bigger focus” trend...

Basically, AI training is learning, and AI inference is doing. Every time you interact with AI—like when you ask ChatGPT a question and it answers—that's inference. The training happened earlier behind the scenes before the AI was released into the wild.

We're hearing a lot about companies starting to focus more on inference these days.

So it's important to recognize that Innodata's solutions aren't just training AIs. They also play vital roles in inference by:

- Providing testing services to ensure models perform like they're supposed to in the real world;
- Optimizing models for faster response times;
- Reducing computing requirements through techniques like “knowledge distillation;”
- Testing models for vulnerabilities that could lead to harmful outputs;
- Continuously monitoring and evaluating outputs for safety;
- Ensuring AI outputs meet regulatory requirements;
- Providing fine tuning to improve inference quality;
- Enhancing performance for particular use cases or industries.

You may have heard of something called “synthetic data” in the context of AI development.

Synthetic data is just what it sounds like—artificially generated data that mimics real data. It's becoming more important in AI development because of:

- Privacy regulations: Restrictions on using personal data for AI training.
- Data scarcity: A lack of sufficient real-world examples to teach the AIs, especially for rare events.

- Bias mitigation: The need to balance and diversify datasets.
- Cost efficiency: Reducing the expense of collecting and annotating real data.
- Edge case coverage: Creating examples of scenarios too rare to collect naturally.

Innodata has positioned itself as a leader in this emerging field of synthetic data, helping its customers overcome data limitations while maintaining privacy and compliance.

Innodata's path to profits

The way the AI market is developing is playing right into Innodata's strengths.

As I noted earlier, Innodata generated revenue of \$170.5 million in 2024, which reflected extremely robust 96% growth from the year before.

The company could certainly maintain that growth rate for the next few years. But I like to lean to the conservative side when making growth forecasts.

I think Innodata can grow revenue at a compound annual rate of at least 60% for the next few years. Using that forecasted growth rate, I calculate a GARP Quotient of 0.2. That's firmly in the "extremely promising" tier we talked about in [The Growth Stock Anomaly special report](#).

In terms of profitability, Innodata reached that milestone in 2024, reporting net income of \$28.7 million.

That said, I don't expect much in the way of earnings for the next few years as the company invests heavily in growth. But it's an added safety net to know that Innodata can operate profitably today if it needs to.

Action to take: Buy one-third of a position in Innodata (INOD) at current market prices. Wait for future guidance to buy more and to set a stop.

(Important: We're recommending you only purchase one-third of your desired allocation to Innodata today. That's because Innodata's stock will likely take a big hit in this market if one or more of the big tech companies comes out with reduced AI investment expectations.

Plus, Innodata's stock is up about 200% since mid-September, so we wouldn't be surprised to see a pullback in this market. It makes sense to keep some powder dry and look to acquire more Innodata shares at lower prices. We plan to buy in two more tranches, but if the stock takes a hit in the near term, we may end up buying the remaining two-thirds all at once. We'll let you know.)

America's fastest-growing stocks

In this monthly feature, we share data from our proprietary AI about America's fastest-growing stocks.

In the table below, we share 25 of the fastest-growing stocks in America that fit our *Disruption_X* size criteria. As a reminder, the stocks we're targeting have a market cap between \$200 million and \$20 billion.

Future recommendations may come from this list. It also gives us great insight into what's growing fast and—just as important—what's not. A stock riding on pure hype can't make this list because it requires real, bona fide revenue growth.

Note these stocks trade on American exchanges. They aren't necessarily American companies.

We also include their raw GARP Quotient based on their most recent quarterly revenue, to give an idea of how fairly valued they are.

Observations this month:

- **Immersion Corp. (IMMR)** leads the list with massive 3,539% revenue growth, which looks more like a one-time licensing pop than sustainable momentum.
- Aerospace & Defense continues to gain steam. Names like **EHang Holdings Ltd. (EH)** and **Intuitive Machines (LUNR)** reflect rising demand in both defense and advanced air mobility.
- Hardware and chip infrastructure is driving real-world innovation again, with multiple entrants showing triple-digit growth in sectors like semiconductors, power systems, and quantum computing.
- Two *Disruption_X* holdings are on the list: **Credo Technology Group Holding Ltd. (CRDO)** and this month's new pick, Innodata.

A few fast-growing stocks that have our attention:

EHang Holdings Ltd. (EH)—Chinese air taxi. EHang makes autonomous aerial vehicles (AAVs)—essentially drone taxis—and it's already flying commercially in China. With 356% revenue growth and a \$1.46 billion market cap, this is a pure play on the future of urban air mobility. FAA and international approvals will be the next big catalysts.

Reitar Logtech Holdings Ltd. (RITR)—Construction automation. RITR provides high-end logistics and construction engineering systems. Revenue surged 187% over the past year, and the company is still under most investors' radars with a modest \$195 million market cap. GARP is at a reasonable 1.2, suggesting growth is coming at a fair price.

IonQ Inc. (IONQ)—Quantum computing. IonQ builds quantum processors that could soon outperform even today's fastest supercomputers. Revenue grew 95%, but what's more notable is the company's backlog and commercial traction. With an increasing enterprise interest, IonQ could become a foundational layer of the next computing era.

This month's list showcases a cross-section of real innovation—air taxis, quantum chips, clean infrastructure, and entertainment tech. The strongest growth is increasingly coming from companies building physical and digital infrastructure for the next decade.

Fastest-growing stocks in America							
Rank	Ticker	Name	Description	Mkt cap (million)	Stock Price	Revenue Growth	GARP Quotient
1	IMMR	Immersion Corporation	Haptic Software	\$253	\$7.82	3,539%	0
2	EH	EHang Holdings Limited	Chinese Air Taxi	\$1,460	\$23.01	288%	0.5
3	MVIS	MicroVision, Inc.	Lidar Sensors	\$364	\$1.62	278%	149.5
4	ICG	Intchains Group Limited	Crypto Chipmaker	\$237	\$3.95	243%	0.4
5	RITR	Reitar Logtech	Construction Automation	\$195	\$2.76	187%	1.2
6	CLSK	CleanSpark, Inc.	Bitcoin Miner	\$2,470	\$8.79	119%	4.5
7	ZK	ZEEKR	Chinese EV Maker	\$6,520	\$25.49	47%	0.1
8	RIVN	Rivian Automotive, Inc.	Electric Truckmaker	\$13,780	\$12.19	12%	17.2
9	LUNR	Intuitive Machines, Inc.	Lunar Spacecraft	\$1,390	\$9.22	134%	5.3
10	IONQ	IonQ, Inc.	Quantum Computing	\$5,820	\$26.12	95%	188.3
11	ADSE	ADS-TEC Energy PLC	EV Power Systems	\$781	\$14.50	130%	2
12	CRDO	Credo Technology	Data Center Chips	\$8,580	\$50.53	99%	24.9
13	SMCI	Super Micro Computer, Inc.	AI Server Hardware	\$24,760	\$41.72	125%	0.6
14	MAX	MediaAlpha, Inc.	Insurance Marketing	\$664	\$9.91	123%	0.5
15	HEPS	D-Market Elektronik	Turkish E-Commerce	\$977	\$3.04	97%	0
16	BYRN	Byrna Technologies Inc.	Non-Lethal Weapons	\$427	\$18.87	101%	5.8
17	MTN	Vail Resorts, Inc.	Ski Resort Operator	\$6,050	\$162.08	4%	41
18	MSGE	Madison Square Garden	Entertainment Venues	\$1,630	\$34.10	10%	10.8
19	SOUN	SoundHound AI, Inc.	Voice AI Software	\$4,110	\$10.31	85%	43.1

Fastest-growing stocks in America							
Rank	Ticker	Name	Description	Mkt cap (million)	Stock Price	Revenue Growth	GARP Quotient
20	AEVA	Aeva Technologies, Inc.	Lidar Chips	\$259	\$4.74	110%	20.2
21	LCID	Lucid Group, Inc.	Luxury EV Maker	\$7,490	\$2.47	36%	27.6
22	INOD	Innodata Inc.	AI Data Services	\$1,440	\$46.12	96%	4.7
23	ASPN	Aspen Aerogels, Inc.	Aerogel Insulation	\$580	\$7.07	90%	2.4
24	IDCC	InterDigital, Inc.	Wireless IP Licensing	\$5,690	\$221.60	58%	7.4
25	VFS	VinFast Auto Ltd.	Vietnamese EV Maker	\$7,900	\$3.38	71%	0

For the most up-to-date guidance on our current positions, [click here](#) to view the portfolio online.

Chris Wood Stephen McBride

Chris Wood and Stephen McBride

Keep up with RiskHedge on the go.

Download the App



Scan it with your Phone

[READ IMPORTANT DISCLOSURES HERE.](#)

YOUR USE OF THESE MATERIALS IS SUBJECT TO THE TERMS OF THESE DISCLOSURES.

Copyright © 2025 RiskHedge. All Rights Reserved